

Device Drivers

Slide 1 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science

Overview

- Device drivers implement a uniform interface to varied hardware devices
 - ◆ This allows the rest of the kernel to be hardware-independent
- Device drivers convert high-level function calls, like `read()`, into specific hardware commands
 - ◆ These commands are typically reads/writes to specific addresses in the physical address space that are assigned to each device
 - ◆ Long-latency commands signal completion via interrupts
- Note that device drivers may be layered
 - ◆ eg, a command may first be handled by a device driver that manages bus transactions, then a device driver that controls the specific device

Slide 2 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science

Memory-Mapped I/O

- Regions of the physical address space are assigned to each device on the bus
 - ◆ These addresses are usually only accessible by the kernel
 - ◆ Multiple devices may share a single controller

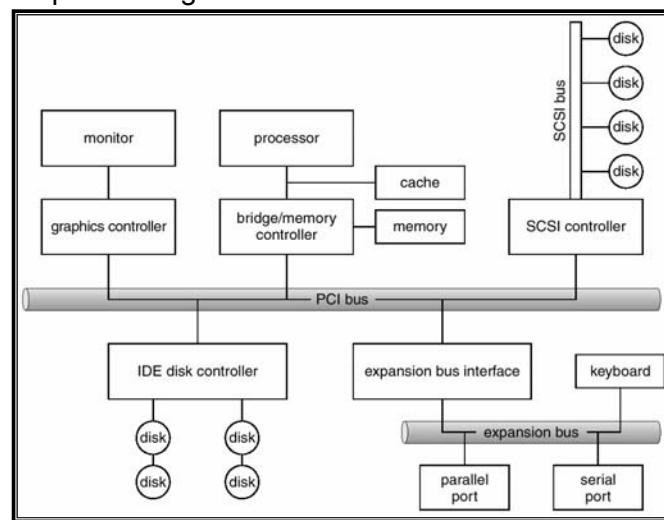
Slide 3 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science

Typical Bus Organization

- Example bus organization:



Slide 4 (© Kevin Skadron, 2004)

Source: Shuberschatz, Galvin and Gagne ©2002,
<http://www.cs.yale.edu/homes/avios/book/03/index.html>



U.Va.
Department of Computer Science

Memory-Mapped I/O, cont.

I/O address range (hexadecimal)	device
000-00F	DMA controller
020-021	interrupt controller
040-043	timer
200-20F	game controller
2F8-2FF	serial port (secondary)
320-32F	hard-disk controller
378-37F	parallel port
3D0-3DF	graphics controller
3F0-3F7	diskette-drive controller
3F8-3FF	serial port (primary)

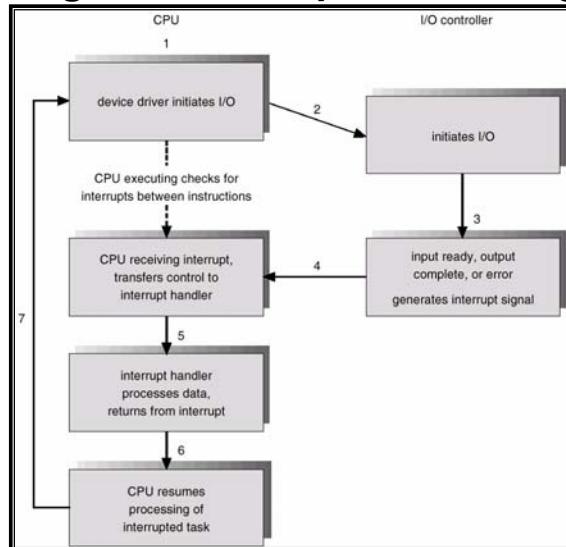
Source: Shiverschatz, Galvin and Gagne ©2002,
<http://cs-www.cs.yale.edu/home/avios/bookos/index.html>

Slide 5 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science

Stages of Interrupt Processing



Source: Shiverschatz, Galvin and Gagne ©2002,
<http://cs-www.cs.yale.edu/home/avios/bookos/index.html>

Slide 6 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science

Sample Interrupts

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19-31	(Intel reserved, do not use)
32-255	maskable interrupts

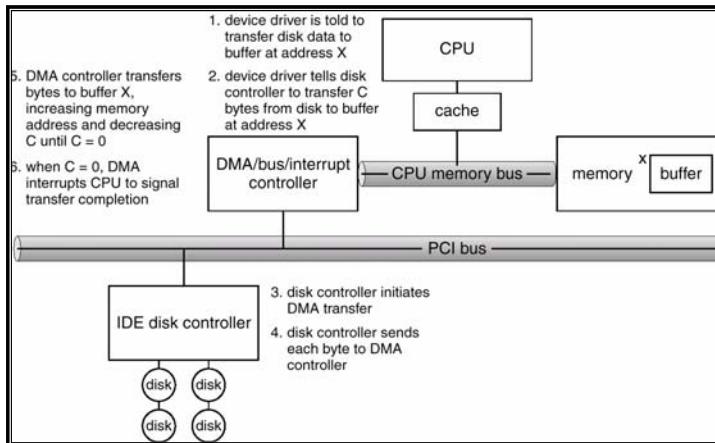
Slide 7 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science

Source: Silberschatz, Galvin and Gagne ©2002,
<http://cs-www.cs.yale.edu/homes/avios/book/osc/index.html>

Stages of DMA



Slide 8 (© Kevin Skadron, 2004)

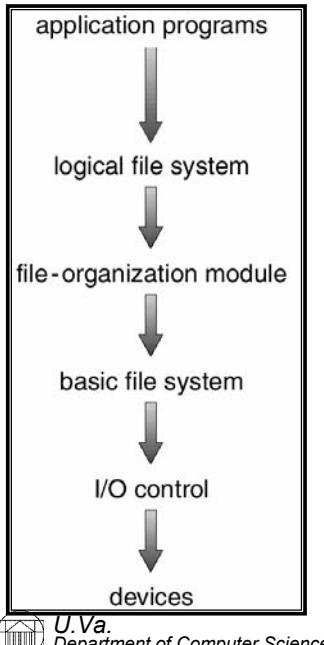


U.Va.
Department of Computer Science

Source: Silberschatz, Galvin and Gagne ©2002,
<http://cs-www.cs.yale.edu/homes/avios/book/osc/index.html>

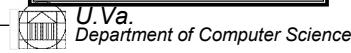
Device Infrastructure

- Hierarchy of abstraction
 - ◆ Device-dependence is confined to I/O control and device driver layers

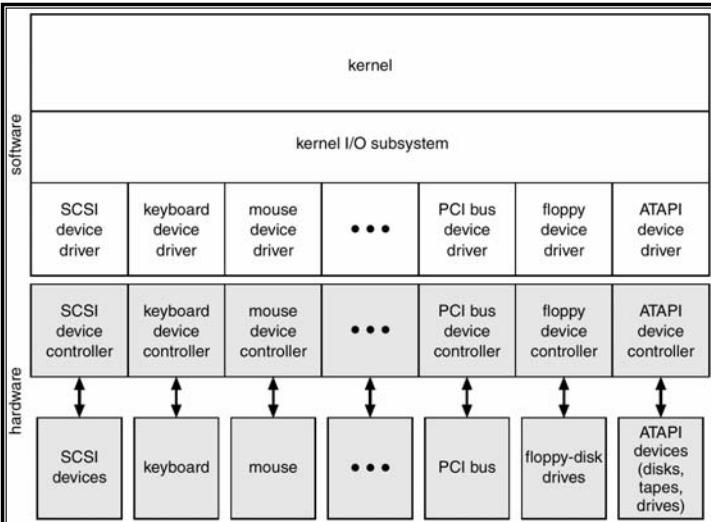


Source: Silberschatz, Galvin and Gagne ©2002,
<http://cs-www.cs.yale.edu/home/avios/book/osc/index.html>

Slide 9 (© Kevin Skadron, 2004)

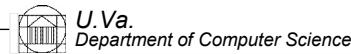


Kernel I/O Structure



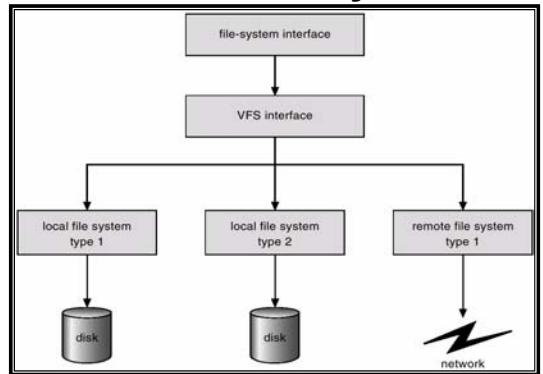
Source: Silberschatz, Galvin and Gagne ©2002,
<http://cs-www.cs.yale.edu/home/avios/book/osc/index.html>

Slide 10 (© Kevin Skadron, 2004)



More abstraction: Virtual File System

- Even the file system has several layers



Source: Silberschatz, Galvin and Gagne ©2012,
<http://cs-www.cs.yale.edu/classes/avios/book/osc/index.html>

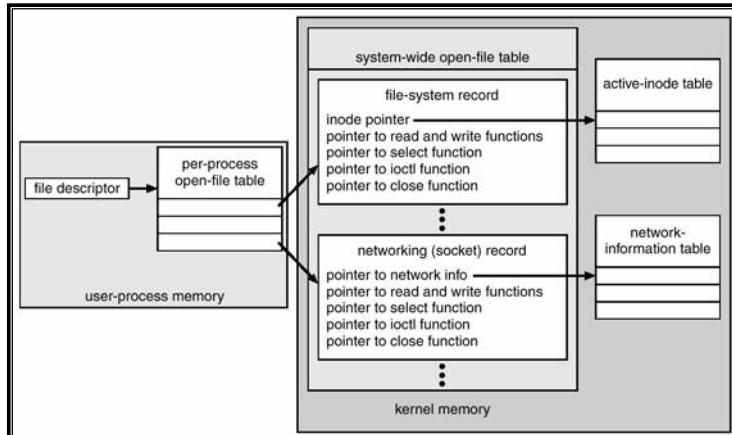
- Virtual File System (VFS) layer – distinguishes local files from remote ones, and local files are further distinguished according to their file-system types.*
 - ◆ The VFS activates file-system-specific operations to handle local requests according to their file-system types.
 - ◆ Calls the NFS protocol procedures for remote requests.

Slide 11 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science

Unix File-Access Control Structure

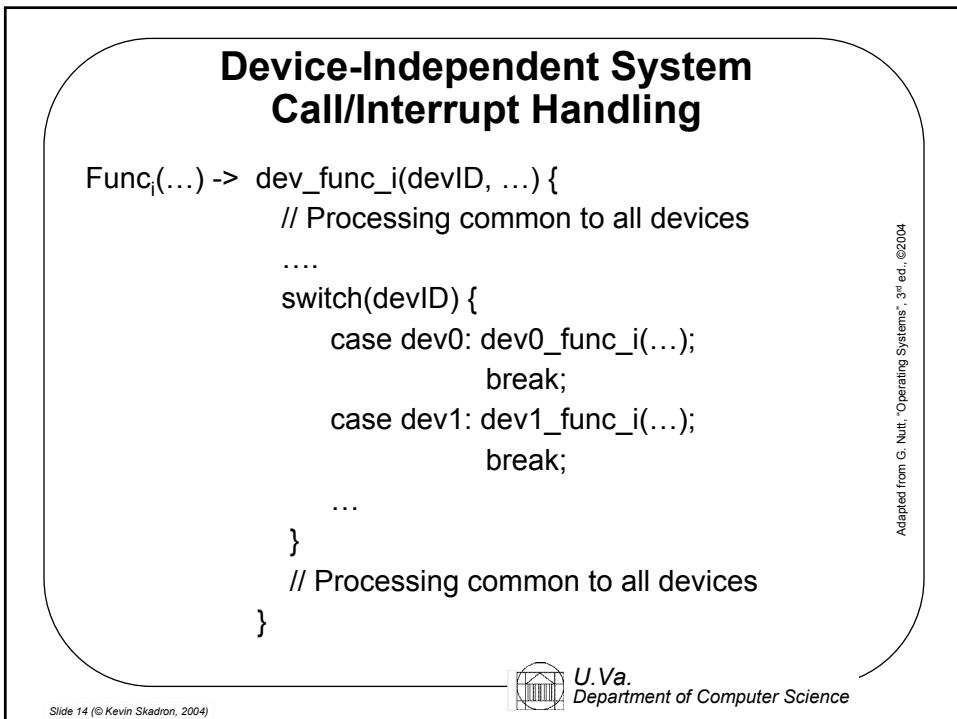
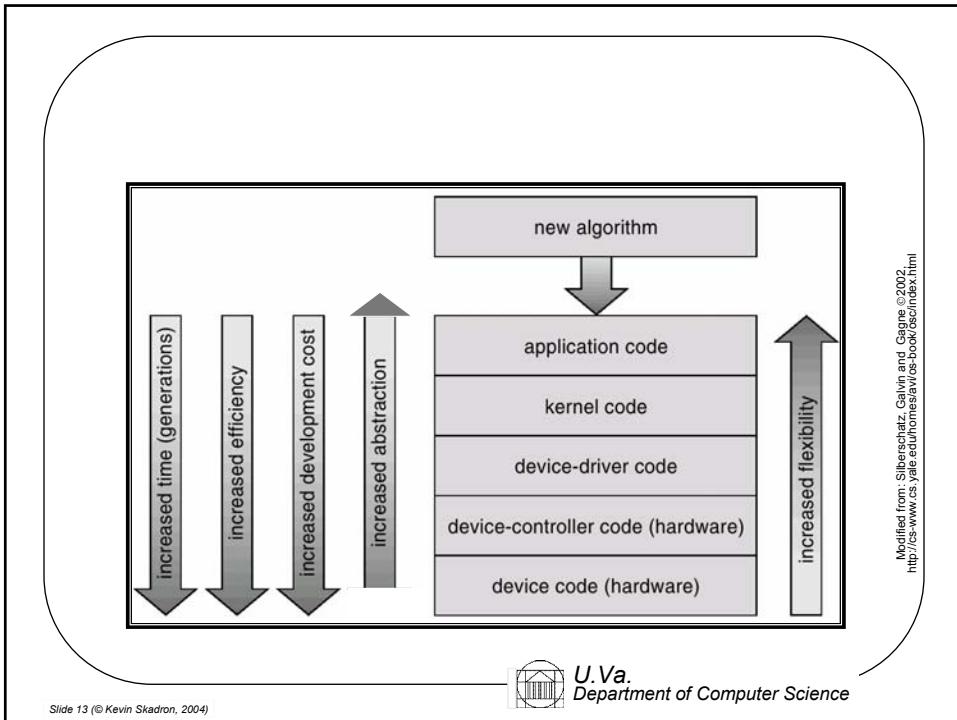


Source: Silberschatz, Galvin and Gagne ©2002,
<http://cs-www.cs.yale.edu/classes/avios/book/osc/index.html>

Slide 12 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science



Example: Linux

```
struct file_operations{  
    loff_t (*llseek) (struct file *, loff_t, int);  
    ssize_t (*read) (struct file *, char *, size_t, loff_t *)  
    ...  
}  
  
foo_init() {  
    ...  
    struct file_operations foo_fops = {  
        NULL;           /* llseek – not used */  
        foo_read;      /* read function */  
    }  
}
```

Slide 15 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science

Example Device Driver Internals

- Interrupt handler for keyboard
 - ◆ (see <http://www.ibiblio.org/navigator-bin/navigator.cgi?drivers/char/keyboard.c>)
 - On screen

Slide 16 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science

Review: Stages in Processing an I/O Call

- Application makes call to C library function, eg fscanf()
- C library translates this into a system call, eg read()
- Execute syscall instruction
- Trap vector processes file name, determines which device is being accessed
 - ◆ This may entail additional disk accesses
- VFS consulted
- Appropriate local file system is consulted (or NFS access)
- File name -> inode -> block number
- Read(block num) call to device driver
 - ◆ Device driver should check input validity, avoid buffer overflow
 - ◆ Device driver checks device status; if busy, request is queued
- Read command issued over bus
 - ◆ This may actually entail a series of back-and-forth bus transactions
- Hardware controller performs read
- Interrupt generated when read completes
- Interrupt handler calls appropriate handler func in device driver

Slide 17 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science

Interrupt Processing

- Interrupt hardware saves some registers, indexes kernel's interrupt vector
- Kernel can ignore interrupt or proceed
- If proceeding, kernel typically temporarily suspends interrupts
- Kernel saves registers and sets up context for the interrupt service procedure
 - ◆ If a user process was running, this requires a context switch, TLB/MMU setup, and creation of a stack
- Run the “top half” of the interrupt service procedure
- Re-enable interrupts
- Finish interrupt servicing
- This may make a process runnable
- Choose which process to run
- Call kernel scheduler

Slide 18 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science

Re-entrant Code

- Note that an interrupt may be received at any time unless interrupts are disabled
 - ◆ While kernel is running
 - ◆ While device driver is handling a system call
 - ◆ While device driver is handling an interrupt
- Disabling interrupts briefly allows kernel/drivers to do minimal work necessary to maintain a coherent state, then interrupts can be re-enabled

Slide 19 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science

More Complexity: Hot-Swappable Devices

- Some devices can be installed or removed while computer runs, eg USB devices
- Either action generates an interrupt
- On an install, kernel must find, (install if necessary), and initialize the device driver
- On a removal, kernel must shut down the device driver and gracefully deal with pending requests
 - ◆ eg, Waiting processes shouldn't hang

Slide 20 (© Kevin Skadron, 2004)



U.Va.
Department of Computer Science